

# Classify different data sets

## Basic includes

```
In [1]: # Using pandas to load the csv file
import pandas as pd

import numpy as np
import matplotlib.pyplot as plt

from keras import models
from keras import layers
from keras import callbacks
from keras import optimizers
from keras.utils import to_categorical

# reuters and fashin mnist data set from keras
from keras.datasets import reuters
from keras.datasets import fashion_mnist

# needed to preprocess text
from keras.preprocessing.text import Tokenizer
```

Using TensorFlow backend.

## Classify the Fashion Mnist


---

```
In [2]: (fashion_train_data, fashion_train_labels), (fashion_test_data, fashion_test_labels) = \
        fashion_mnist.load_data()

print(fashion_train_data.shape)
print(fashion_test_data.shape)
test_index = 10
plt.title("Label: " + str(fashion_train_labels[test_index]))
plt.imshow(fashion_train_data[test_index], cmap="gray")

(60000, 28, 28)
(10000, 28, 28)
```

Out[2]: <matplotlib.image.AxesImage at 0x132872438>



## Standardizing images

standardizing data images and creating one hot labels

### TO DO: Preprocess the data

1. Normalize the input data set
2. Perform one hot encoding
3. Create a train, test, and validation set

```
In [3]: fashion_train_data = fashion_train_data.reshape((60000, 28 * 28))
fashion_train_data = fashion_train_data.astype('float32') / 255

fashion_test_data = fashion_test_data.reshape((10000, 28 * 28))
fashion_test_data = fashion_test_data.astype('float32') / 255

# one hot encoding
fashion_one_hot_labels = to_categorical(fashion_train_labels)
fashion_test_one_hot_labels = to_categorical(fashion_test_labels)

#creating validation set for first 10000 elements
fashion_validation_data = fashion_train_data[:10000]
fashion_validation_labels = fashion_one_hot_labels[:10000]

#creating input set
x_data = fashion_train_data[10000:]
y_data = fashion_one_hot_labels[10000:]
print(x_data.shape)
print(y_data.shape)

(50000, 784)
(50000, 10)
```

**TO DO: Define and train a network, then plot the accuracy of the training, validation, and testing**

1. Use a validation set
2. Propose and train a network
3. Print the history of the training
4. Evaluate with a test set

```
In [4]: #building keras model
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_dim=784))
model.add(layers.Dropout(0.4))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.4))
model.add(layers.Dense(len(fashion_one_hot_labels[0]), activation='softmax'))
model.summary()

# included the early stopping which monitors the validation loss
early_stop = callbacks.EarlyStopping(monitor='val_loss', patience=4)
model.compile(loss='categorical_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	200960
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290

=====  
 Total params: 235,146  
 Trainable params: 235,146  
 Non-trainable params: 0  
 =====

```
In [5]: history = model.fit(x_data, y_data,
                           batch_size=512,
                           epochs=40,
                           validation_data=(fashion_validation_data, fashion_validation_labels),
                           callbacks=[early_stop],
                           verbose=2)
```

Train on 50000 samples, validate on 10000 samples

Epoch 1/40

- 2s - loss: 0.6291 - acc: 0.6657 - val\_loss: 0.4280 - val\_acc: 0.7825

Epoch 2/40

- 2s - loss: 0.4106 - acc: 0.7980 - val\_loss: 0.3610 - val\_acc: 0.8192

Epoch 3/40

- 1s - loss: 0.3633 - acc: 0.8208 - val\_loss: 0.3200 - val\_acc: 0.838

```
3
Epoch 4/40
- 2s - loss: 0.3384 - acc: 0.8318 - val_loss: 0.3055 - val_acc: 0.845
6
Epoch 5/40
- 1s - loss: 0.3211 - acc: 0.8400 - val_loss: 0.3114 - val_acc: 0.844
4
Epoch 6/40
- 2s - loss: 0.3102 - acc: 0.8457 - val_loss: 0.2901 - val_acc: 0.854
6
Epoch 7/40
- 1s - loss: 0.3003 - acc: 0.8495 - val_loss: 0.2853 - val_acc: 0.856
5
Epoch 8/40
- 2s - loss: 0.2950 - acc: 0.8525 - val_loss: 0.2885 - val_acc: 0.855
6
Epoch 9/40
- 2s - loss: 0.2883 - acc: 0.8558 - val_loss: 0.2675 - val_acc: 0.866
0
Epoch 10/40
- 2s - loss: 0.2850 - acc: 0.8579 - val_loss: 0.2670 - val_acc: 0.866
5
Epoch 11/40
- 2s - loss: 0.2795 - acc: 0.8604 - val_loss: 0.2687 - val_acc: 0.864
5
Epoch 12/40
- 1s - loss: 0.2736 - acc: 0.8631 - val_loss: 0.2750 - val_acc: 0.862
9
Epoch 13/40
- 1s - loss: 0.2698 - acc: 0.8643 - val_loss: 0.2606 - val_acc: 0.868
7
Epoch 14/40
- 1s - loss: 0.2671 - acc: 0.8667 - val_loss: 0.2611 - val_acc: 0.869
2
Epoch 15/40
- 1s - loss: 0.2636 - acc: 0.8682 - val_loss: 0.2662 - val_acc: 0.866
1
Epoch 16/40
- 1s - loss: 0.2602 - acc: 0.8700 - val_loss: 0.2561 - val_acc: 0.870
9
Epoch 17/40
- 1s - loss: 0.2592 - acc: 0.8703 - val_loss: 0.2820 - val_acc: 0.858
0
Epoch 18/40
- 2s - loss: 0.2574 - acc: 0.8713 - val_loss: 0.2502 - val_acc: 0.874
0
Epoch 19/40
- 2s - loss: 0.2520 - acc: 0.8733 - val_loss: 0.2553 - val_acc: 0.869
8
Epoch 20/40
```

```
- 2s - loss: 0.2513 - acc: 0.8744 - val_loss: 0.2463 - val_acc: 0.876
2
Epoch 21/40
- 2s - loss: 0.2515 - acc: 0.8746 - val_loss: 0.2649 - val_acc: 0.866
8
Epoch 22/40
- 2s - loss: 0.2471 - acc: 0.8763 - val_loss: 0.2414 - val_acc: 0.878
7
Epoch 23/40
- 2s - loss: 0.2499 - acc: 0.8745 - val_loss: 0.2679 - val_acc: 0.865
7
Epoch 24/40
- 2s - loss: 0.2479 - acc: 0.8758 - val_loss: 0.2438 - val_acc: 0.876
0
Epoch 25/40
- 2s - loss: 0.2459 - acc: 0.8773 - val_loss: 0.2463 - val_acc: 0.875
5
Epoch 26/40
- 1s - loss: 0.2423 - acc: 0.8790 - val_loss: 0.2499 - val_acc: 0.875
0
```

```

In [9]: #evaluating the model with the test data
results = model.evaluate(fashion_test_data, fashion_test_one_hot_labels)
print(results)

history_dict = history.history
print(history_dict.keys())

#creating list variables for plotting validation
acc = history_dict['acc']
val_acc = history_dict['val_acc']
loss = history_dict['loss']
val_loss = history_dict['val_loss']
epochs = range(1, len(acc) + 1)

#plotting validation and training loss

plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()


plt.show()

```

```

10000/10000 [=====] - 0s 33us/step
[0.2633243887126446, 0.8671]
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```




```
In [8]: #plotting validation and train accuracy

plt.clf()

plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



## Classifying newswires

Build a network to classify Reuters newswires into 46 different mutually-exclusive topics.

### Load and review the data



```
In [3]: (reuters_train_data, reuters_train_labels),(reuters_test_data, reuters_t

print(reuters_train_data.shape)
print(reuters_train_labels.shape)
print(reuters_train_data[0])
print(reuters_train_labels[0])
print(set(reuters_train_labels))

(8982,)
(8982,)
[1, 2, 2, 8, 43, 10, 447, 5, 25, 207, 270, 5, 3095, 111, 16, 369, 186,
 90, 67, 7, 89, 5, 19, 102, 6, 19, 124, 15, 90, 67, 84, 22, 482, 26, 7,
 48, 4, 49, 8, 864, 39, 209, 154, 6, 151, 6, 83, 11, 15, 22, 155, 11, 1
 5, 7, 48, 9, 4579, 1005, 504, 6, 258, 6, 272, 11, 15, 22, 134, 44, 11,
 15, 16, 8, 197, 1245, 90, 67, 52, 29, 209, 30, 32, 132, 6, 109, 15, 17
 , 12]
3
{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19,
 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37
 , 38, 39, 40, 41, 42, 43, 44, 45}
```

Load the word index to decode the train data.

```
In [4]: word_index = reuters.get_word_index()

reverse_index = dict([(value+3, key) for (key, value) in word_index.items()])

reverse_index[0] = "<PAD>"
reverse_index[1] = "<START>"
reverse_index[2] = "<UNKNOWN>" # unknown
reverse_index[3] = "<UNUSED>"

decoded_review = ' '.join([reverse_index.get(i, '?') for i in reuters_train_data[0]])

print(decoded_review)
```

<START> <UNKNOWN> <UNKNOWN> said as a result of its december acquisiti  
on of space co it expects earnings per share in 1987 of 1 15 to 1 30 d  
lrs per share up from 70 cts in 1986 the company said pretax net shoul  
d rise to nine to 10 mln dlrs from six mln dlrs in 1986 and rental ope  
ration revenues to 19 to 22 mln dlrs from 12 5 mln dlrs it said cash f  
low per share this year should be 2 50 to three dlrs reuter 3

**TO DO: Preprocess the data**

1. Normalize the input data set
2. Perform one hot encoding
3. Create a train, test, and validation set

```
In [5]: tokenizer = Tokenizer(reuters_train_data.shape[0])
reuters_train_data_token = tokenizer.sequences_to_matrix(reuters_train_data)
reuters_test_data_token = tokenizer.sequences_to_matrix(reuters_test_data)
print(reuters_train_data_token.shape)
print(reuters_test_data_token.shape)

# One-hot encoding the output
reuters_one_hot_train_labels = to_categorical(reuters_train_labels)
reuters_one_hot_test_labels = to_categorical(reuters_test_labels)
print(reuters_one_hot_train_labels.shape)
print(reuters_one_hot_test_labels.shape)

# Creating a validation set with the first 10000 reviews
reuters_validation_data = reuters_train_data_token[:3000]
reuters_validation_labels = reuters_one_hot_train_labels[:3000]
print(reuters_validation_data.shape)
print(reuters_validation_labels.shape)
# Creating the input set for the
x_data = reuters_train_data_token[3000:]
y_data = reuters_one_hot_train_labels[3000:]
print("x:", x_data.shape)
print("y:", y_data.shape)

(8982, 8982)
(2246, 8982)
(8982, 46)
(2246, 46)
(3000, 8982)
(3000, 46)
x: (5982, 8982)
y: (5982, 46)
```

```
In [6]: #building keras model
model1 = models.Sequential()
model1.add(layers.Dense(256, activation='relu', input_dim=len(reuters_tr
model1.add(layers.Dropout(0.5))
# model1.add(layers.Dropout(0.3))
model1.add(layers.Dense(len(reuters_one_hot_test_labels[0]), activation=
model1.summary()

# included the early stopping which monitors the validation loss
early_stop = callbacks.EarlyStopping(monitor='val_loss', patience=8)
model1.compile(loss='categorical_hinge',
               optimizer='adadelta',
               metrics=[ 'accuracy' ])
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	2299648
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 46)	11822
Total params: 2,311,470		
Trainable params: 2,311,470		
Non-trainable params: 0		

```
In [7]: history = model1.fit(x_data, y_data,
                             batch_size=256,
                             epochs=50,
                             validation_data=(reuters_validation_data, reuters_validation_l
                             callbacks=[early_stop],
                             verbose=2)
```

Train on 5982 samples, validate on 3000 samples

Epoch 1/50

- 3s - loss: 1.0027 - acc: 0.3009 - val\_loss: 0.9933 - val\_acc: 0.4930

Epoch 2/50

- 3s - loss: 0.8711 - acc: 0.4689 - val\_loss: 0.7503 - val\_acc: 0.5180

Epoch 3/50

- 3s - loss: 0.6960 - acc: 0.5401 - val\_loss: 0.6380 - val\_acc: 0.5770

Epoch 4/50

- 3s - loss: 0.5893 - acc: 0.5834 - val\_loss: 0.5755 - val\_acc: 0.6307

Epoch 5/50

```
- 3s - loss: 0.5547 - acc: 0.6133 - val_loss: 0.5611 - val_acc: 0.659
3
Epoch 6/50
- 3s - loss: 0.5399 - acc: 0.6456 - val_loss: 0.5531 - val_acc: 0.665
7
Epoch 7/50
- 2s - loss: 0.5283 - acc: 0.6603 - val_loss: 0.5491 - val_acc: 0.678
0
Epoch 8/50
- 2s - loss: 0.5190 - acc: 0.6854 - val_loss: 0.5431 - val_acc: 0.687
0
Epoch 9/50
- 2s - loss: 0.5104 - acc: 0.6954 - val_loss: 0.5345 - val_acc: 0.696
7
Epoch 10/50
- 2s - loss: 0.4962 - acc: 0.7098 - val_loss: 0.5242 - val_acc: 0.696
3
Epoch 11/50
- 2s - loss: 0.4735 - acc: 0.7150 - val_loss: 0.5063 - val_acc: 0.696
0
Epoch 12/50
- 3s - loss: 0.4532 - acc: 0.7245 - val_loss: 0.4900 - val_acc: 0.704
3
Epoch 13/50
- 2s - loss: 0.4314 - acc: 0.7379 - val_loss: 0.4811 - val_acc: 0.718
3
Epoch 14/50
- 2s - loss: 0.4175 - acc: 0.7441 - val_loss: 0.4720 - val_acc: 0.720
0
Epoch 15/50
- 2s - loss: 0.4000 - acc: 0.7546 - val_loss: 0.4666 - val_acc: 0.727
0
Epoch 16/50
- 2s - loss: 0.3871 - acc: 0.7646 - val_loss: 0.4596 - val_acc: 0.733
3
Epoch 17/50
- 2s - loss: 0.3763 - acc: 0.7753 - val_loss: 0.4561 - val_acc: 0.734
3
Epoch 18/50
- 2s - loss: 0.3657 - acc: 0.7823 - val_loss: 0.4540 - val_acc: 0.738
3
Epoch 19/50
- 2s - loss: 0.3546 - acc: 0.7909 - val_loss: 0.4488 - val_acc: 0.744
3
Epoch 20/50
- 2s - loss: 0.3450 - acc: 0.7991 - val_loss: 0.4450 - val_acc: 0.745
0
Epoch 21/50
- 2s - loss: 0.3354 - acc: 0.8061 - val_loss: 0.4407 - val_acc: 0.750
3
```

```
Epoch 22/50
- 3s - loss: 0.3269 - acc: 0.8104 - val_loss: 0.4363 - val_acc: 0.757
3
Epoch 23/50
- 2s - loss: 0.3177 - acc: 0.8183 - val_loss: 0.4318 - val_acc: 0.760
3
Epoch 24/50
- 2s - loss: 0.3098 - acc: 0.8206 - val_loss: 0.4293 - val_acc: 0.759
0
Epoch 25/50
- 2s - loss: 0.3005 - acc: 0.8255 - val_loss: 0.4270 - val_acc: 0.765
3
Epoch 26/50
- 2s - loss: 0.2925 - acc: 0.8290 - val_loss: 0.4233 - val_acc: 0.767
3
Epoch 27/50
- 3s - loss: 0.2850 - acc: 0.8307 - val_loss: 0.4198 - val_acc: 0.767
7
Epoch 28/50
- 2s - loss: 0.2796 - acc: 0.8368 - val_loss: 0.4203 - val_acc: 0.768
3
Epoch 29/50
- 2s - loss: 0.2741 - acc: 0.8372 - val_loss: 0.4157 - val_acc: 0.770
0
Epoch 30/50
- 2s - loss: 0.2679 - acc: 0.8397 - val_loss: 0.4153 - val_acc: 0.771
3
Epoch 31/50
- 2s - loss: 0.2615 - acc: 0.8432 - val_loss: 0.4110 - val_acc: 0.774
7
Epoch 32/50
- 3s - loss: 0.2554 - acc: 0.8460 - val_loss: 0.4105 - val_acc: 0.778
0
Epoch 33/50
- 2s - loss: 0.2517 - acc: 0.8482 - val_loss: 0.4077 - val_acc: 0.778
0
Epoch 34/50
- 3s - loss: 0.2489 - acc: 0.8501 - val_loss: 0.4063 - val_acc: 0.779
3
Epoch 35/50
- 2s - loss: 0.2415 - acc: 0.8526 - val_loss: 0.4032 - val_acc: 0.782
3
Epoch 36/50
- 2s - loss: 0.2373 - acc: 0.8547 - val_loss: 0.4041 - val_acc: 0.779
7
Epoch 37/50
- 3s - loss: 0.2341 - acc: 0.8571 - val_loss: 0.4054 - val_acc: 0.780
0
Epoch 38/50
- 2s - loss: 0.2315 - acc: 0.8587 - val_loss: 0.4046 - val_acc: 0.780
```

```

7
Epoch 39/50
  - 2s - loss: 0.2283 - acc: 0.8616 - val_loss: 0.4007 - val_acc: 0.783
3
Epoch 40/50
  - 2s - loss: 0.2247 - acc: 0.8639 - val_loss: 0.3988 - val_acc: 0.783
7
Epoch 41/50
  - 2s - loss: 0.2206 - acc: 0.8641 - val_loss: 0.4006 - val_acc: 0.783
7
Epoch 42/50
  - 3s - loss: 0.2178 - acc: 0.8699 - val_loss: 0.3962 - val_acc: 0.786
3
Epoch 43/50
  - 2s - loss: 0.2163 - acc: 0.8701 - val_loss: 0.3959 - val_acc: 0.787
3
Epoch 44/50
  - 2s - loss: 0.2128 - acc: 0.8719 - val_loss: 0.3997 - val_acc: 0.786
3
Epoch 45/50
  - 2s - loss: 0.2111 - acc: 0.8723 - val_loss: 0.3931 - val_acc: 0.788
0
Epoch 46/50
  - 2s - loss: 0.2078 - acc: 0.8801 - val_loss: 0.3941 - val_acc: 0.791
0
Epoch 47/50
  - 2s - loss: 0.2050 - acc: 0.8801 - val_loss: 0.3935 - val_acc: 0.789
0
Epoch 48/50
  - 2s - loss: 0.2028 - acc: 0.8828 - val_loss: 0.3912 - val_acc: 0.791
0
Epoch 49/50
  - 2s - loss: 0.1995 - acc: 0.8838 - val_loss: 0.3939 - val_acc: 0.792
7
Epoch 50/50
  - 2s - loss: 0.1982 - acc: 0.8858 - val_loss: 0.3910 - val_acc: 0.794
3

```

```

In [8]: #evaluating the model with the test data
print(reuters_test_data.shape)
print(reuters_one_hot_test_labels.shape)
results = model1.evaluate(reuters_test_data_token, reuters_one_hot_test_
print(results)

history_dict = history.history
print(history_dict.keys())

#creating list variables for plotting validation
acc = history_dict['acc']
val_acc = history_dict['val acc']

```

```

loss = history_dict['loss']
val_loss = history_dict['val_loss']
epochs = range(1, len(acc) + 1)


#plotting validation and training loss

plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

(2246,)
(2246, 46)
2246/2246 [=====] - 0s 185us/step
[0.39702498440432527, 0.7800534283700843]
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```




```
In [9]: #plotting validation and train accuracy

plt.clf()

plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



**TO DO: Define and train a network, then plot the accuracy of the training, validation, and testing**

1. Use a validation set
2. Propose and train a network
3. Print the history of the training
4. Evaluate with a test set



# Predicting Student Admissions

---

Predict student admissions based on three pieces of data:

- GRE Scores
- GPA Scores
- Class rank

## Load and visualize the data

```
In [2]: student_data = pd.read_csv("student_data.csv")  
print(student_data)
```

	admit	gre	gpa	rank
0	0	380.0	3.61	3.0
1	1	660.0	3.67	3.0
2	1	800.0	4.00	1.0
3	1	640.0	3.19	4.0
4	0	520.0	2.93	4.0
5	1	760.0	3.00	2.0
6	1	560.0	2.98	1.0
7	0	400.0	3.08	2.0
8	1	540.0	3.39	3.0
9	0	700.0	3.92	2.0
10	0	800.0	4.00	4.0
11	0	440.0	3.22	1.0
12	1	760.0	4.00	1.0
13	0	700.0	3.08	2.0
14	1	700.0	4.00	1.0
15	0	480.0	3.44	3.0
16	0	780.0	3.87	4.0
17	0	360.0	2.56	3.0
18	0	800.0	3.75	2.0
19	1	540.0	3.81	1.0
20	0	500.0	3.17	3.0
21	1	660.0	3.63	2.0
22	0	600.0	2.82	4.0
23	0	680.0	3.19	4.0
24	1	760.0	3.35	2.0
25	1	800.0	3.66	1.0
26	1	620.0	3.61	1.0
27	1	520.0	3.74	4.0
28	1	780.0	3.22	2.0
29	0	520.0	3.29	1.0

```
..      ...      ...      ...      ...
370      1  540.0  3.77  2.0
371      1  680.0  3.76  3.0


372      1  680.0  2.42  1.0
373      1  620.0  3.37  1.0
374      0  560.0  3.78  2.0
375      0  560.0  3.49  4.0
376      0  620.0  3.63  2.0
377      1  800.0  4.00  2.0
378      0  640.0  3.12  3.0
379      0  540.0  2.70  2.0
380      0  700.0  3.65  2.0
381      1  540.0  3.49  2.0
382      0  540.0  3.51  2.0
383      0  660.0  4.00  1.0
384      1  480.0  2.62  2.0
385      0  420.0  3.02  1.0
386      1  740.0  3.86  2.0
387      0  580.0  3.36  2.0
388      0  640.0  3.17  2.0
389      0  640.0  3.51  2.0
390      1  800.0  3.05  2.0
391      1  660.0  3.88  2.0
392      1  600.0  3.38  3.0
393      1  620.0  3.75  2.0
394      1  460.0  3.99  3.0
395      0  620.0  4.00  2.0
396      0  560.0  3.04  3.0
397      0  460.0  2.63  2.0
398      0  700.0  3.65  2.0
399      0  600.0  3.89  3.0
```

```
[400 rows x 4 columns]
```

Plot of the GRE and the GPA from the data.

```
In [3]: X = np.array(student_data[["gre", "gpa"]])
y = np.array(student_data["admit"])
admitted = X[np.argwhere(y==1)]
rejected = X[np.argwhere(y==0)]
plt.scatter([s[0][0] for s in rejected], [s[0][1] for s in rejected], s=50, color='red')
plt.scatter([s[0][0] for s in admitted], [s[0][1] for s in admitted], s=50, color='cyan')
plt.xlabel('Test (GRE)')
plt.ylabel('Grades (GPA)')

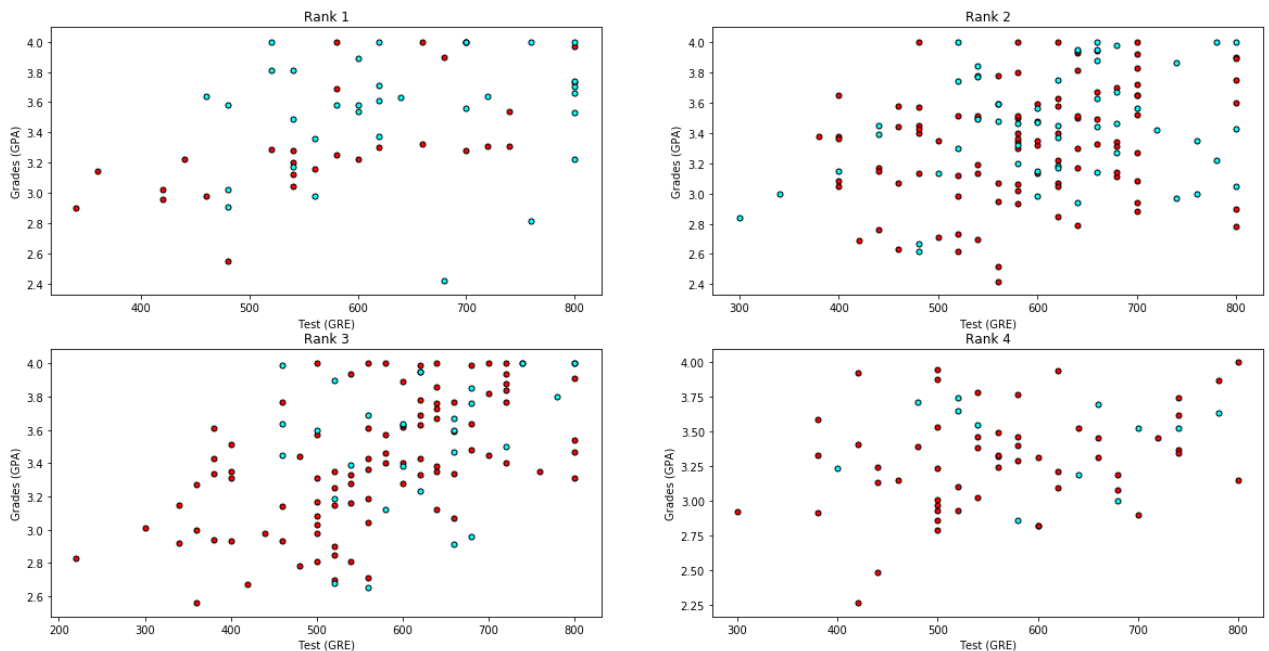
plt.show()
```



Plot of the data by class rank.

```
In [4]: f, plots = plt.subplots(2, 2, figsize=(20,10))
plots = [plot for sublist in plots for plot in sublist]

for idx, plot in enumerate(plots):
    data_rank = student_data[student_data["rank"]==idx+1]
    plot.set_title("Rank " + str(idx+1))
    X = np.array(data_rank[["gre", "gpa"]])
    y = np.array(data_rank["admit"])
    admitted = X[np.argwhere(y==1)]
    rejected = X[np.argwhere(y==0)]
    plot.scatter([s[0][0] for s in rejected], [s[0][1] for s in rejected])
    plot.scatter([s[0][0] for s in admitted], [s[0][1] for s in admitted])
    plot.set_xlabel('Test (GRE)')
    plot.set_ylabel('Grades (GPA)')
```



### TO DO: Preprocess the data

1. Normalize the input data set
2. Perform one hot encoding
3. Create a train, test, and validation set

```

In [5]: # admit    gre    gpa    rank
student_data.fillna(value= 0, inplace= True)
gpa = np.array(student_data["gpa"])
gpa = (gpa - np.nanmean(gpa))/np.nanstd(gpa)
gre = np.array(student_data["gre"])
gre = (gre - np.nanmean(gre))/np.nanstd(gre)
rank = np.array(student_data["rank"])
rank = to_categorical(rank)
standarizedStudentData = np.zeros((len(gre), 2))
standarizedStudentData[:,0] = gre
standarizedStudentData[:,1] = gpa
studentDataFinal = np.zeros((len(gre), 7))
for i, (studentData, cat) in enumerate(zip(standarizedStudentData, rank))
    studentDataFinal[i] = np.concatenate((studentData, cat))

admit = np.array(student_data["admit"])
studentOneHotLabels = to_categorical(admit)

print(studentDataFinal.shape)
print(studentOneHotLabels.shape)

studentTrainData = studentDataFinal[:300]
studentTrainLabels = studentOneHotLabels[:300]
studentTestData = studentDataFinal[300:]
studentTestLabels = studentOneHotLabels[300:]

(400, 7)
(400, 2)

```

Type Markdown and LaTeX:  $\alpha^2$

```
In [6]: #building keras model
model2 = models.Sequential()
model2.add(layers.Dense(32, activation='relu', input_dim=7))
# model1.add(layers.Dropout(0.4))
model2.add(layers.Dense(32, activation='relu'))
# model1.add(layers.Dropout(0.3))
model2.add(layers.Dense(2, activation='sigmoid'))
model2.summary()

# included the early stopping which monitors the validation loss
early_stop = callbacks.EarlyStopping(monitor='val_loss', patience=8)
model2.compile(loss='kullback_leibler_divergence',
               optimizer='nadam',
               metrics=[ 'accuracy' ])
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_1 (Dense)	(None, 32)	256
dense_2 (Dense)	(None, 32)	1056
dense_3 (Dense)	(None, 2)	66
=====	=====	=====
Total params: 1,378		
Trainable params: 1,378		
Non-trainable params: 0		
=====		

```
In [7]: history = model2.fit(studentTrainData, studentTrainLabels,
                             batch_size=512,
                             epochs=200,
                             validation_split=0.2,
                             verbose=2)

Train on 240 samples, validate on 60 samples
Epoch 1/200
- 0s - loss: 0.6918 - acc: 0.3958 - val_loss: 0.6550 - val_acc: 0.5833
Epoch 2/200
- 0s - loss: 0.6635 - acc: 0.4917 - val_loss: 0.6377 - val_acc: 0.6333
Epoch 3/200
- 0s - loss: 0.6432 - acc: 0.5625 - val_loss: 0.6217 - val_acc: 0.6667
Epoch 4/200
- 0s - loss: 0.6244 - acc: 0.6208 - val_loss: 0.6058 - val_acc: 0.7167
Epoch 5/200
- 0s - loss: 0.6059 - acc: 0.6542 - val_loss: 0.5900 - val_acc: 0.7000
Epoch 6/200
- 0s - loss: 0.5873 - acc: 0.6792 - val_loss: 0.5739 - val_acc: 0.6500
- 1s - loss: 0.5700 - acc: 0.6958 - val_loss: 0.5575 - val_acc: 0.6833
```

```
In [8]: #evaluating the model with the test data
print(studentTestData.shape)
print(studentTestLabels.shape)
results = model2.evaluate(studentTestData, studentTestLabels)
print(results)

history_dict = history.history
print(history_dict.keys())


#creating list variables for plotting validation
acc = history_dict['acc']
val_acc = history_dict['val_acc']
loss = history_dict['loss']
val_loss = history_dict['val_loss']
epochs = range(1, len(acc) + 1)

#plotting validation and training loss

plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
```

```
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
  
plt.show()
```

```
(100, 7)  
(100, 2)  
100/100 [=====] - 0s 41us/step  
[0.0012386897555552424, 0.65]  
dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```






```
In [9]: #plotting validation and train accuracy

plt.clf()

plt.plot(epochs, acc, 'r', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



**TO DO: Define and train a network, then plot the accuracy of the training, validation, and testing**

1. Use a validation set
2. Propose and train a network
3. Print the history of the training
4. Evaluate with a test set